


For each unit of work, the router decides which model runs it. It computes the tier the work needs from checkable facts about the task, then picks a model with capacity in that tier. Every decision is deterministic, explainable, logged, and recomputed at each attempt. It works for hives on API keys, subscriptions, local models, or a mix.



Deciding the tier

Tiers: local · light · standard · strong. The tier is the highest one any rule requires:

- F1** Kind: digest local; work, review light; plan, triage standard; consult strong.
 - F2** Only a goal, no plan: standard, and suggest splitting off a plan task.
 - F3** Touches many components: standard.
 - F4** No check at all: standard.
 - F5** Costly or hard to undo, without an independent check: strong.
 - F6** Many tasks blocked on it, without an independent check: strong.
 - F7** Reviews: at least the author's tier.
 - F8** Two failed checks: one tier up (at strong, a consult).
- Facts, not verdicts.** They come from the task and the record: is there a plan, is there a check the worker can't edit, how much does it touch, what did earlier attempts do.
 - Cheaper means better-specified.** The creator's hint can only raise the tier. To go cheaper, add an independent check, write the plan first, or split the task.
 - Scorer.** A small fixed model estimates *unknown* facts only, through the same rules. Shadow mode first; live once replay shows it helps.
 - Learning.** Replay computes success per kind, facts and tier; a person approves each table change.

Choosing the route

- RT1** Operator override wins (recorded).
- RT2** Only qualified routes that fit: pool has capacity, cost fits budget, context and tools fit. Unknown counts as not fitting.
- RT3** The tier is a **floor**. Nothing fits now → wait; nothing could ever fit → no_route (a table problem). Never a silent downgrade.
- RT4** Reviews use a different model family.
- RT5** After a failure, act on its class (right); a reassignment excludes the last route.
- RT6** Within the tier: listed order, lowest cost, or most headroom.
- RT7** After an interruption, truncation or missing info, keep the route.

Core ideas

- Route:** a named, pinned model config; any change makes a new version, so switches are visible. Each passes a **canary** on past tasks, with a recorded lesson; qualification is bound to the pin.
- Pool:** capacity with limits: **metered** (API money), **subscription** (usage windows), **local** (concurrency).
- One route per attempt** or per episode (a route service for agents with their own loop), re-routed at restarts and checkpoints; **scripts, not models**, for polling, parsing and tests.

Failures: one remedy per class

Class	Remedy
outage	Another route in the same tier, or wait. Never a tier up.
capacity	Rate limit, window or budget hit: another pool, or wait.
truncated	Larger limits, or split the task. Not a tier up.
missing info	Get the information, retry on the same route.
failed check	Retry once with feedback, then a tier up or a consult.
stalled	Supervisor or operator decides; a tier up is one option.
interrupted	Process lost or restarted, not the model's fault: same route.
indeterminate	Unknown whether the call ran: reconcile first, no blind retry.

Patterns the rules produce

- Plan strong, execute light.** A plan task yields a plan and checks; the work tasks are then well specified and checked, so they run light.
- Consult.** A strong model answers one narrow question, with no tools; a light worker does the work.

Gateway: LiteLLM, where a route needs it

- Routes are gateway aliases; agents get a worker key that can only call models, never provider credentials.
- Spend is logged per pool, route, task and attempt; tag budgets cap each period and each attempt.
- The route table is the source of truth; gateway fallbacks and hidden retries are off.

How it fits One Hive

R1 record	route.* summaries (amendment A1): decided, waiting, canary, drift, mode, table; each bound to the router's full, replayable log entry.
R5 review	Rules F7 and RT4.
R6 runtime	Asks the router at each attempt start; launches the harness with the answer (examples/launch.sh).
R9 replay	Replays under another table; evaluates the scorer; proposes tuning.

Rollout, and what "done" means

Built and live in 1-hive: decide() and its rules; usage from providers' own reports and the gateway; canaries, drift checks, the scorer (shadow); a gateway for self-hosted models and API keys, with budgets per period and per attempt and a worker key; routing on the hive record. **Adopting:** docs/ADOPTING.md: route table → launcher in **fixed** mode (the baseline) → canaries → **live** → tuning. **Done:** more accepted tasks per unit of capacity than the baseline, with no drop in review pass rate.

Not in v1: the scorer setting tiers directly, automatic tuning, a spend ledger, mid-attempt switches.